

Capstone Project Documentation

Food Delivery Analytics

End-to-End Data Engineering on Azure Databricks

Platform: Azure Databricks (Runtime 17.3, Photon Engine)

Cloud Provider: Microsoft Azure

Storage: Azure Data Lake Storage Gen2 (ADLS)

Data Format: Delta Lake

Architecture: Medallion Architecture (Bronze → Silver → Gold)

Serving Layer: Power BI via External Tables in hive_metastore

Notebook	Layer	Purpose
01_Bronze_Ingestion	Bronze	Raw CSV ingestion from ADLS → Delta folders
02_Silver_Transformation	Silver	Data cleaning, CDC merge, SCD Type 2 processing
03_Gold_Modeling	Gold	Star schema modeling, KPI aggregation, Z-Order optimization

1. Executive Summary

This capstone project implements a **production-grade data engineering pipeline** for a Food Delivery Analytics Platform using the **Medallion Architecture** on **Azure Databricks**. The platform processes transactional food delivery data — including orders, users, and restaurants — through three progressively refined layers of data quality.

The pipeline begins by ingesting raw CSV files from Azure Data Lake Storage (ADLS Gen2) into the **Bronze layer** as immutable Delta folders, preserving the original data with full lineage metadata. The **Silver layer** then applies business-critical transformations including string standardization, Change Data Capture (CDC) merges for real-time order status updates, and Slowly Changing Dimension (SCD) Type 2 processing for user and restaurant dimensions. Finally, the **Gold layer** models the cleansed data into a **Star Schema** (fact and dimension tables) and pre-computes business KPIs for analytical consumption.

All data assets are stored as Delta Lake folders on ADLS Gen2 and registered as **External Tables** in Databricks `hive_metastore`, enabling seamless integration with **Power BI** for executive dashboards and operational reporting.

Key Highlights

- **4 raw data sources** ingested (orders, orders_cdc, users, restaurants)
- **CDC merge logic** using Delta Lake's `MERGE INTO` for real-time order status tracking
- **SCD Type 2** dimension processing for historical change tracking
- **Star Schema** with 2 dimension tables and 1 fact table
- **3 pre-computed KPI tables** for dashboard performance optimization
- **Z-Order optimization** on the fact table for time-series query acceleration
- **Photon engine** enabled for 2-5x query performance improvement
- **End-to-end auditability** with ingestion timestamps and source file tracking

2. Business Problem Statement

2.1 Problem Context

A growing food delivery company generates thousands of orders daily across multiple cities and restaurant partners. The business leadership requires **real-time visibility** into operational performance, customer behavior, and restaurant partner metrics to make data-driven decisions. However, the existing data infrastructure faces several critical challenges:

- **Fragmented Data Sources:** Order transactions, user profiles, and restaurant metadata exist as disconnected CSV files with no unified data model.
- **Stale Order Status:** Order statuses change frequently (Ordered → Preparing → Delivered / Cancelled), but the base system only captures snapshot data, missing incremental updates.
- **Dimension Drift:** User and restaurant attributes (city, cuisine type, ratings) change over time with no mechanism to track historical states.
- **No Analytical Layer:** Raw data cannot be directly consumed by BI tools — it requires cleaning, deduplication, and schema modeling.
- **Performance Bottlenecks:** Ad-hoc queries on flat files are slow and resource-intensive, hindering real-time dashboard performance.

2.2 Business Objectives

#	Objective	Measured By
1	Unified Data Platform — Consolidate all food delivery data into a single, governed data lake	All 4 sources ingested and queryable
2	Real-Time Order Tracking — Apply CDC to reflect the latest order statuses	CDC merge successfully updates order status
3	Historical Dimension Tracking — Implement SCD Type 2 for user and restaurant changes	Audit columns (<code>_processed_at</code>) track changes

4	Analytical Star Schema — Model data for efficient BI consumption	Fact + Dimension tables in Gold layer
5	Pre-Computed KPIs — Accelerate dashboard load times with pre-aggregated metrics	3 KPI tables: city revenue, restaurant performance, daily trends
6	Scalable Architecture — Design for future growth with Delta Lake and Medallion pattern	Idempotent, rerunnable pipeline with overwrite semantics


2.3 Target Stakeholders

- **Business Executives:** Revenue dashboards, city-level performance insights
- **Operations Team:** Daily order volume trends, delivery status monitoring
- **Restaurant Partners Team:** Restaurant performance rankings, cuisine-wise analysis
- **Data Engineering Team:** Pipeline health, data quality monitoring, schema evolution

3. Technical Architecture (Medallion Structure)

3.1 Architecture Overview

The pipeline follows the **Medallion Architecture** — a multi-layered data design pattern that progressively refines data from raw ingestion to analytics-ready outputs.



3.2 Technology Stack

Component	Technology	Purpose
Cloud Platform	Microsoft Azure	Infrastructure hosting
Storage	ADLS Gen2 (satactivity)	Scalable data lake storage
Compute	Azure Databricks (DBR 17.3)	Distributed data processing
Engine	Photon	Vectorized query acceleration
Data Format	Delta Lake	ACID transactions, time travel, schema enforcement
Language	PySpark + SQL	Data transformation and modeling
Metastore	hive_metastore	External table registration
BI Tool	Power BI	Dashboard visualization
Auth	ADLS Access Key	Storage account authentication

3.3 Layer 1: Bronze (Raw Ingestion)

Notebook: 01_Bronze_Ingestion

The Bronze layer performs a faithful, 1:1 ingestion of raw CSV files into Delta format with zero business logic applied. It serves as the **immutable audit layer** — the single source of truth for raw data.

Process Flow

1. **Authentication** — Configures Spark with ADLS access key for the satactivity storage account

2. **Cleanup** — Recursively deletes the existing bronze/ directory for idempotent reruns
3. **Ingestion Loop** — Iterates over 4 source files with a reusable helper function:
 - a. Reads CSV with header and schema inference
 - b. Adds metadata columns: `_ingestion_timestamp` (current time) and `_source_file_name` (source path)
 - c. Writes to Delta format with overwrite mode
4. **Verification** — Lists the resulting Delta folders to confirm successful ingestion

Source Files → Bronze Folders

Source CSV	Bronze Folder	Description
orders.csv	bronze/orders/	Base order transactions
orders_cdc.csv	bronze/orders_cdc/	Incremental order status changes
users_scd.csv	bronze/users/	User dimension with SCD attributes
restaurants_scd.csv	bronze/restaurants/	Restaurant dimension with SCD attributes

Key Design Decisions

- **Delta Format from Day 1** — Even at the Bronze level, Delta provides ACID transactions, schema enforcement, and time travel
- **Metadata Enrichment** — `_ingestion_timestamp` and `_source_file_name` enable full data lineage tracking
- **Idempotent Design** — Full cleanup + overwrite ensures the pipeline can be safely rerun without duplicates
- **No Schema Enforcement** — `inferSchema=true` allows flexible ingestion of evolving source schemas

3.4 Layer 2: Silver (Cleaned & Conformed)

Notebook: 02_Silver_Transformation

The Silver layer applies **data quality rules, CDC merge logic, and dimension standardization** to produce clean, conformed datasets.

Process Flow

1. **Data Cleaning** — A reusable `clean_dataframe()` function trims whitespace and applies `initcap()` to all string columns
2. **Orders CDC Merge** — The most complex transformation:
 - a. Loads both `bronze/orders` (base) and `bronze/orders_cdc` (incremental changes)
 - b. Deduplicates CDC records using a window function (`ROW_NUMBER` partitioned by `order_id`, ordered by `updated_at DESC`) to keep only the latest change per order
 - c. Uses Delta Lake's **MERGE INTO** (`DeltaTable.merge()`) to:
 - i. **Update** existing orders with new status values
 - ii. **Insert** new orders not yet in Silver
 - d. Handles first-run initialization: if Silver orders don't exist, creates the initial table then runs the merge
3. **Dimension Processing** — A reusable `process_dimension()` function:
 - a. Reads Bronze dimension data
 - b. Applies string cleaning
 - c. Adds `_processed_at` audit column
 - d. Writes to Silver with overwrite mode

Transformations Applied

Transformation	Applied To	Technique
String Trimming	All string columns	<code>trim()</code>
Initial Caps	All string columns	<code>initcap()</code>
CDC Deduplication	<code>orders_cdc</code>	<code>ROW_NUMBER()</code> window function
CDC Merge	<code>orders</code>	Delta <code>MERGE INTO</code> (upsert)
Audit Timestamps	All tables	<code>current_timestamp()</code>
Schema Evolution	Orders merge	<code>mergeSchema=true</code> option

Key Design Decisions

- **CDC via Delta MERGE** — Provides exactly-once semantics for order status updates, preventing duplicates and ensuring consistency
- **Window-Based Deduplication** — Ensures only the most recent CDC event per order is applied, handling out-of-order arrivals
- **Reusable Functions** — `clean_dataframe()` and `process_dimension()` enforce consistent cleaning across all entities
- **First-Run Safety** — Try/except pattern handles both initial loads and subsequent incremental runs gracefully

3.5 Layer 3: Gold (Analytical Model)

Notebook: 03_Gold_Modeling

The Gold layer models Silver data into a **Star Schema** and pre-computes **KPI aggregates** for Power BI consumption.

Star Schema Design



Process Flow

- Dimension Tables** — Select relevant reporting attributes from Silver:
 - dim_users: user_id, user_name, city
 - dim_restaurants: restaurant_id, restaurant_name, cuisine, rating
- Fact Table** — Joins Silver orders with both dimensions via LEFT JOINs, denormalizing city, restaurant_name, and cuisine into the fact for query performance
- KPI Aggregates** — Three pre-computed metric tables:
 - kpi_revenue_by_city: Total revenue and order count per city
 - kpi_restaurant_performance: Revenue, average rating, and order count per restaurant
 - kpi_daily_trends: Daily revenue and order count time series
- Table Registration** — Each Gold table is saved as a Delta folder and registered as an External Table in hive_metastore with gold_ prefix

5. **Optimization** — Z-Order indexing on fact_orders by order_timestamp for time-series query acceleration

Gold Tables Inventory

Table Name	Type	Description	Key Columns
gold_dim_users	Dimension	User attributes for reporting	user_id, user_name, city
gold_dim_restaurants	Dimension	Restaurant attributes	restaurant_id, restaurant_name, cuisine, rating
gold_fact_orders	Fact	Denormalized order transactions	order_id, user_id, restaurant_id, total_amount, status, order_timestamp
gold_kpi_revenue_by_city	KPI Aggregate	City-level revenue summary	city, total_revenue, total_orders
gold_kpi_restaurant_performance	KPI Aggregate	Restaurant rankings	restaurant_name, revenue, avg_rating, order_count
gold_kpi_daily_trends	KPI Aggregate	Daily order trends	order_date, daily_revenue, daily_orders

4. Data Quality and Failure Handling

4.1 Data Quality Measures by Layer

Data quality is enforced progressively across each layer, following the principle of "**clean as you go**" — each layer adds more stringent quality checks.

Bronze Layer — Ingestion Integrity

Quality Measure	Implementation	Purpose
Schema Inference	inferSchema=true on CSV read	Automatically detects correct data types (int, string, timestamp)
Header Validation	header=true on CSV read	Ensures column names are read from the file header, not auto-generated
Lineage Tracking	_ingestion_timestamp, _source_file_name columns	Full audit trail: when was data ingested and from which source file
Immutability	Delta format with overwrite semantics	Bronze data serves as the immutable raw audit layer
Completeness Check	dbutils.fs.ls(bronze_path) verification	Confirms all 4 expected folders are created after ingestion

Silver Layer — Transformation Quality

Quality Measure	Implementation	Purpose
String Standardization	trim() + initcap() on all string columns	Eliminates whitespace inconsistencies and casing variations
CDC Deduplication	ROW_NUMBER() window function on order_id ordered by updated_at DESC	Ensures only the latest status change per order is applied

Upsert Integrity	Delta MERGE INTO with match/not-match conditions	Prevents duplicate orders; ensures exactly-once processing
Schema Evolution	mergeSchema=true on initial write	Handles cases where CDC introduces new columns
Default Value Handling	lit("Ordered") fallback for missing status column	Ensures dimensional completeness even with sparse source data
Audit Timestamps	_last_updated (orders), _processed_at (dimensions)	Tracks when each record was last modified in Silver

Gold Layer — Analytical Quality

Quality Measure	Implementation	Purpose
Referential Integrity	LEFT JOINS between fact and dimensions	Ensures all orders appear in the fact table, even if dimension data is missing
Selective Projection	Explicit .select() on dimension columns	Only relevant reporting columns pass through; no internal/audit columns leak
Aggregate Validation	orderBy(desc) on KPIs	Results are pre-sorted for immediate dashboard consumption
Table Registration	DROP TABLE IF EXISTS + CREATE TABLE pattern	Idempotent registration — safe to rerun without leftover metadata

4.2 Failure Handling Strategy

The pipeline implements **defensive programming** at every stage to ensure failures are caught, logged, and recoverable.

Bronze Layer — Ingestion Failures

Ingestion Loop with Try/Except:

```
for file, folder in sources:
    try:
        ingest_to_bronze_folder()
    except Exception as e:
        print("❌ Failed: {e}")
        (continues to next file)
```

- **Per-Source Isolation** — Each source file is ingested independently. A failure in `orders_cdc.csv` does NOT block `users_scd.csv`
- **Error Logging** — Failed ingestions print the exception message for debugging
- **Cleanup Safety** — The cleanup step uses try/except to handle cases where the bronze directory doesn't exist yet
- **Idempotent Reruns** — Full directory cleanup before ingestion means the pipeline can be safely retried

Silver Layer — Transformation Failures

CDC Merge with First-Run Detection:

```
try:
    DeltaTable.forPath(silver/orders)
    → MERGE INTO (upsert)
except:
    → Initialize table first
    → Then run MERGE
```

- **First-Run vs. Incremental** — The try/except pattern auto-detects whether Silver orders exist. On first run, it initializes; on subsequent runs, it merges
- **Missing Column Handling** — If status column is absent in Bronze orders, it is auto-populated with "Ordered" default
- **Schema Flexibility** — mergeSchema=true allows new columns from CDC to be absorbed without pipeline failure

Gold Layer — Modeling Failures

- **DROP + CREATE Pattern** — External tables are dropped before re-creation, preventing "table already exists" errors
- **LEFT JOIN Strategy** — Using LEFT (not INNER) joins ensures no data loss even when dimension records are missing
- **Overwrite Mode** — All Gold writes use mode("overwrite"), making the pipeline fully idempotent

4.3 Data Quality Summary Matrix

Layer	Focus	Key Technique	Failure Recovery
Bronze	Completeness & Lineage	Metadata columns, schema inference	Per-source try/except, cleanup + retry
Silver	Accuracy & Consistency	CDC merge, string standardization	First-run detection, default values
Gold	Integrity & Performance	Star schema, Z-Order optimization	DROP + CREATE idempotency, LEFT joins

5. Final Outputs & Serving Layer

5.1 Output Architecture

All Gold-layer outputs follow a dual-registration pattern:

1. **Physical Storage** — Delta folders on ADLS Gen2 at <abfss://sat-activity@satactivity.dfs.core.windows.net/gold/>
2. **Logical Registration** — External Tables in hive_metastore with gold_ prefix

This separation of physical storage from logical access enables:

- **Direct file access** for Spark workloads (reading Delta paths directly)
- **SQL access** for BI tools (querying registered tables via JDBC/ODBC)
- **Independent lifecycle** — storage and metadata can be managed separately

5.2 Registered External Tables

External Table Name	ADLS Path	Table Type	Row Scope
gold_dim_users	gold/dim_users/	Dimension	All active users
gold_dim_restaurants	gold/dim_restaurants/	Dimension	All restaurants
gold_fact_orders	gold/fact_orders/	Fact (Denormalized)	All orders with joined dimensions
gold_kpi_revenue_by_city	gold/kpi_revenue_by_city/	KPI Aggregate	One row per city
gold_kpi_restaurant_performance	gold/kpi_restaurant_performance/	KPI Aggregate	One row per restaurant

gold_kpi_daily_trends	gold/kpi_daily_trends/	KPI Aggregate	One row per day
-----------------------	------------------------	------------------	--------------------

5.3 KPI Definitions

Three pre-computed KPI tables eliminate the need for expensive runtime aggregations in Power BI:

KPI 1: Revenue by City (gold_kpi_revenue_by_city)

Metric	Expression	Business Question
total_revenue	SUM(total_amount) grouped by city	Which cities generate the most revenue?
total_orders	COUNT(order_id) grouped by city	Which cities have the highest order volume?

KPI 2: Restaurant Performance (gold_kpi_restaurant_performance)

Metric	Expression	Business Question
revenue	SUM(total_amount) grouped by restaurant_name, cuisine	Which restaurants are top earners?
avg_rating	AVG(rating) grouped by restaurant_name, cuisine	What is the average quality score per restaurant?
order_count	COUNT(order_id) grouped by restaurant_name, cuisine	Which restaurants have the most orders?

KPI 3: Daily Trends (gold_kpi_daily_trends)

Metric	Expression	Business Question
--------	------------	-------------------

daily_revenue	SUM(total_amount) grouped by date_trunc('day', order_timestamp)	What is the daily revenue trend?
daily_orders	COUNT(order_id) grouped by date_trunc('day', order_timestamp)	How is order volume trending over time?

5.4 Performance Optimization

Z-Order Indexing

```
OPTIMIZE gold_fact_orders ZORDER BY (order_timestamp)
```

Z-Ordering co-locates related data in the same set of files, dramatically improving query performance for time-based filters. This is critical for:

- **Time-series dashboards** filtering by date ranges
- **Trend analysis** queries scanning order_timestamp
- **Partition elimination** in Power BI date slicers


Photon Engine

The cluster runs **Databricks Photon** — a vectorized query engine written in C++ that provides:

- **2-5x faster** aggregation and join operations
- **Native Delta optimization** for scan and merge operations
- **Reduced compute costs** through more efficient resource utilization

5.5 Power BI Integration

The Gold tables are designed for seamless Power BI consumption:



Connection Method: Databricks JDBC/ODBC connector via Partner Connect or manual configuration

Recommended Dashboard Pages:

Dashboard Page	Primary Table	Key Visuals
Executive Summary	gold_kpi_revenue_by_city	Revenue map, city comparison bar chart
Restaurant Leaderboard	gold_kpi_restaurant_performance	Ranked table, cuisine breakdown pie chart
Daily Operations	gold_kpi_daily_trends	Line chart (revenue), bar chart (order volume)
Order Drillthrough	gold_fact_orders	Detail table with city/restaurant/status filters

6. Conclusion

6.1 Project Summary

This capstone project successfully demonstrates the design and implementation of a **complete, production-grade data engineering pipeline** on Azure Databricks. By leveraging the Medallion Architecture, the pipeline transforms raw CSV files into a fully modeled, optimized, and BI-ready analytical platform.

Achievements

Deliverable	Status	Details
Bronze Ingestion Pipeline	Complete	4 CSV sources → Delta folders with lineage metadata
Silver CDC Merge	Complete	Delta MERGE for real-time order status updates
Silver Dimension Processing	Complete	String standardization + audit timestamps
Gold Star Schema	Complete	2 dimensions + 1 fact table with denormalized attributes
Gold KPI Tables	Complete	3 pre-computed aggregate tables for dashboards
Z-Order Optimization	Complete	Time-series acceleration on fact_orders
External Table Registration	Complete	6 tables registered in hive_metastore
Power BI Readiness	Complete	Tables queryable via JDBC/ODBC for BI consumption

6.2 Technical Competencies Demonstrated

- **Azure Cloud Services** — ADLS Gen2 configuration, storage authentication, ABFSS protocol
- **Apache Spark / PySpark** — DataFrame API, window functions, aggregations, joins
- **Delta Lake** — ACID transactions, MERGE INTO, schema evolution, time travel, Z-Order
- **Data Modeling** — Star schema design, fact/dimension separation, denormalization strategy
- **CDC Processing** — Change Data Capture using Delta Lake merge operations
- **SCD Implementation** — Slowly Changing Dimension processing with audit columns
- **Pipeline Design** — Idempotent, rerunnable workflows with defensive error handling
- **Performance Tuning** — Z-Order indexing, Photon engine, pre-computed aggregates

6.3 Known Limitations

Limitation	Impact	Mitigation
Static CSV ingestion (no streaming)	Data freshness limited to batch frequency	Future: Auto Loader for streaming ingestion
Single-node cluster (0 workers)	Limited parallelism for large datasets	Scale to multi-node for production
Access key authentication	Less secure than managed identity	Future: Azure Key Vault + Service Principal
No data validation framework	Quality checks are implicit in transformations	Future: Great Expectations or Delta Expectations
rating column not included in fact_orders	KPI query for avg_rating fails in Cell 9	Fix: Include r.rating in the fact table SELECT

6.4 Future Enhancements

1. **Streaming Ingestion** — Replace batch CSV reads with Auto Loader (cloudFiles) for near-real-time Bronze ingestion
2. **Unity Catalog Migration** — Move from hive_metastore external tables to Unity Catalog managed tables for centralized governance
3. **Data Quality Framework** — Implement Delta Expectations or Great Expectations for declarative data quality rules
4. **Orchestration** — Schedule the three notebooks as a Databricks Workflow (Job) with dependency chains: Bronze → Silver → Gold
5. **Secret Management** — Replace hardcoded access keys with Azure Key Vault backed Databricks secret scopes
6. **Incremental Processing** — Replace full-overwrite patterns in Silver/Gold with incremental merge logic for better performance at scale
7. **Monitoring & Alerting** — Add pipeline health monitoring with Databricks SQL Alerts for data freshness and quality metrics
8. **Testing** — Add unit tests for transformation functions and integration tests for end-to-end pipeline validation